



Chapter 3:

Neural networks



Outlines

1. Anatomy of neural networks
2. Introduction to Keras
3. Keras in practice



Anatomy of neural networks

- ❖ Training a neural network revolves around the following objects:
 - *Layers*, which are combined into a *network* (or *model*)
 - The *input data* and corresponding *targets*
 - The *loss function*, which defines the feedback signal used for learning
 - The *optimizer*, which determines how learning proceeds

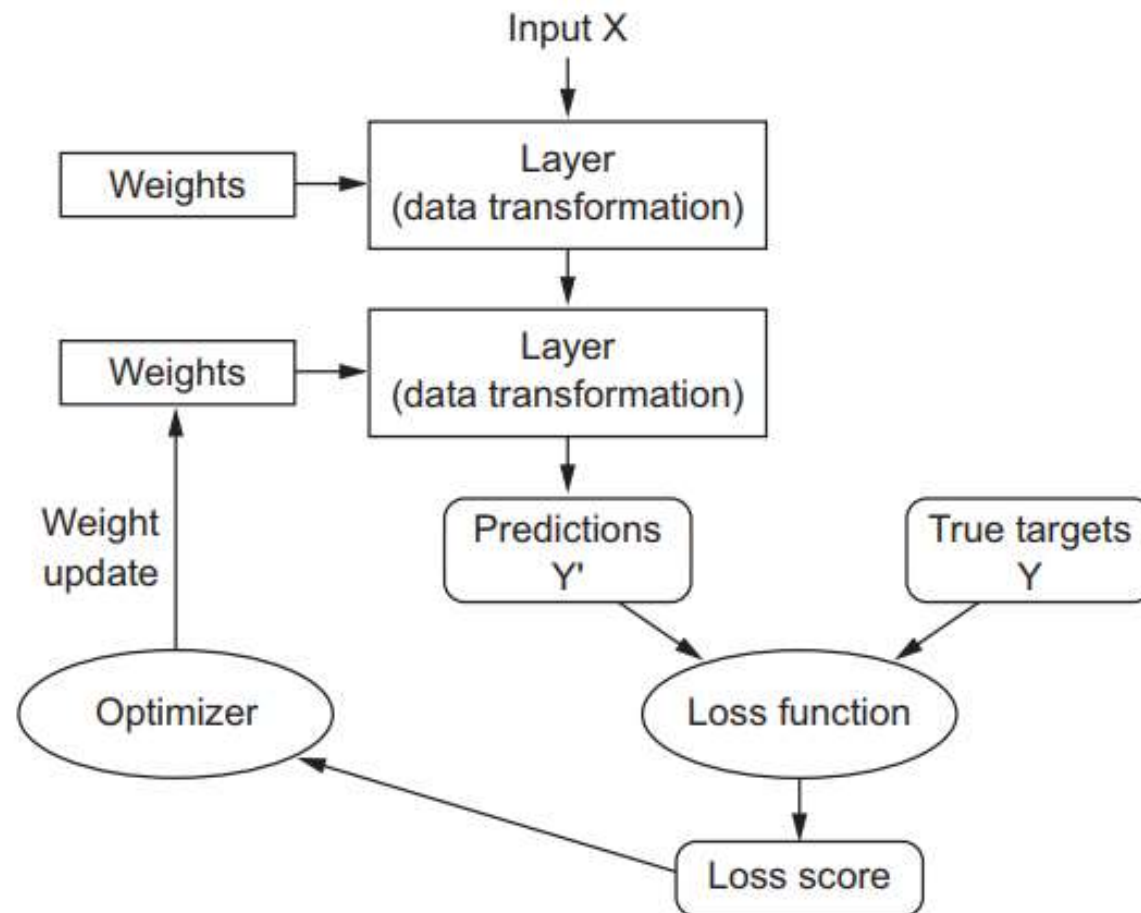


Anatomy of neural networks

- ❖ Training a neural network revolves around the following objects:
 - *Layers*, which are combined into a *network* (or *model*)
 - The *input data* and corresponding *targets*
 - The *loss function*, which defines the feedback signal used for learning
 - The *optimizer*, which determines how learning proceeds

Anatomy of neural networks

- ❖ Relationship between the network, layers, loss function, and optimizer





Anatomy of neural networks

1. The building blocks of deep learning:

- The fundamental data structure in neural networks is the **layer**.
- A layer is a data processing module that takes as input one or more tensors and that outputs one or more tensors.
- **Different types** of **layers** are appropriate for different tensor formats and different types of data processing



Anatomy of neural networks

1. The building blocks of deep learning:

❖ For instance:

➤ Simple vector data:

- ✓ Stored in rank-2 tensors of shape (**samples**, **features**),
- ✓ Often processed by *densely connected* layers, also called *fully connected* or *dense* layers (the Dense class in Keras).

➤ Sequence data:

- ✓ Stored in rank-3 tensors of shape (**samples**, **timesteps**, **features**),
- ✓ These are typically processed by *recurrent* layers, such as an LSTM layer, or 1D convolution layers (**Conv1D**).



Anatomy of neural networks

1. The building blocks of deep learning:

- ❖ For instance (cont):

- Image data:

- ✓ Stored in rank-4 tensors,

- ✓ Usually, 2D convolution layers (**Conv2D**) process this data.



Anatomy of neural networks

1. The building blocks of deep learning:

- ❖ For instance (cont):

- Image data:

- ✓ Stored in rank-4 tensors,

- ✓ Usually, 2D convolution layers (**Conv2D**) process this data.



Anatomy of neural networks

2. Network of layers:

- A deep-learning model is a **directed**, acyclic graph of layers.
- The most common instance is a linear stack of layers, mapping a single input to a single output.
- Some common ones include the following:
 - ✓ Two-branch networks
 - ✓ Multihead networks
 - ✓ Inception blocks



Anatomy of neural networks

2. Network of layers:

- The topology of a network defines a hypothesis space.
- A network topology specifies a series of tensor operations, mapping input data to output data.
- The process the network works on is to find a good set of values for the weight tensors.
- Picking the right network architecture is more an art than a science.
- Practice and practice can help you become a proper neural-network architect.



Anatomy of neural networks

3. Loss functions and optimizers:

- Once the network architecture has been defined, there are two more things to do:
 - *Loss function* (objective function)—The quantity that will be minimized during training. It represents a measure of success for the task at hand.
 - *Optimizer—Determines* how the network will be updated based on the loss function. It implements a specific variant of stochastic gradient descent (SGD).



Anatomy of neural networks

3. Loss functions and optimizers:

- Neural networks with multiple outputs may have multiple loss functions.
- The loss functions must be combined into a single scalar loss for gradient descent.
- Choosing the right objective function is crucial because the network will optimize for it ruthlessly.
- Potentially leading to unintended consequences if the function doesn't fully align with the desired outcome.
- A poorly chosen objective can lead to harm.



Anatomy of neural networks

3. Loss functions and optimizers:

- For common tasks like classification, regression, and sequence prediction, there are standard loss functions:
 - ✓ Binary cross-entropy for two-class classification
 - ✓ Mean squared error for regression.
- Custom loss functions are mainly needed for novel research problems.




Introduction to Keras



- **Keras** is a deep-learning framework for Python that provides a convenient way to define and train almost any kind of deep-learning model.
- **Keras** is a deep-learning framework for Python that provides a convenient way to define and train almost any kind of deep-learning model.
- **Keras** was initially developed for researchers, with the aim of enabling fast experimentation.

Introduction to Keras



- Keras is a high-level API for Neural Networks, written in **Python** that runs on TensorFlow
- **Goal**: go from idea to result with the least possible delay.
- **TensorFlow** is an open source software library for numerical computation using data flow graphs.

Introduction to Keras

➤ **Keras** has the following key features:

- ✓ It allows the same code to run seamlessly on a CPU or GPU.
- ✓ It has a user-friendly API that makes it easy to quickly prototype deep-learning models.
- ✓ It has built-in support for convolutional networks (for computer vision), recurrent networks (for sequence processing), and any combination of both.





Introduction to Keras

- **Keras** is distributed under the permissive **MIT** license
=> it can be freely used in commercial projects.
- **Keras** is compatible with any version of **Python** from 2.7 to 3.6 (as of mid-2017).
- **Keras** has well over 200,000 users.
- **Keras** is used at Google, Netflix, Uber, CERN, Yelp, Square, and hundreds of startups working on a wide range of problems.
- **Keras** is also a popular framework on Kaggle.



Introduction to Keras

1. Keras, Tensorflow, Theano, CNTK:

- Keras is a high-level library for building deep-learning models, relying on specialized tensor libraries for low-level operations.
- Keras supports multiple backend engines, including TensorFlow, Theano, and CNTK, with potential for future expansion.

Introduction to Keras

1. Keras, Tensorflow, Theano, CNTK:

- Keras handles the problem in a modular way; thus, several different backend engines can be plugged seamlessly into Keras.
- There are three existing backend implementations are the TensorFlow backend, the Theano backend, and the Microsoft Cognitive Toolkit (CNTK) backend.



The deep-learning software and hardware stack



Introduction to Keras

1. Keras, Tensorflow, Theano, CNTK:
 - TensorFlow, CNTK, and Theano are major deep-learning platforms, developed by Google, Microsoft, and Université de Montréal, respectively.
 - Keras allows seamless switching between these backends without code changes, making development more flexible.
 - TensorFlow is recommended as the default due to its scalability and widespread use.
 - When running on CPU, TensorFlow is itself wrapping a low-level library for tensor operations called Eigen
 - Keras can run on both CPUs and GPUs, leveraging Eigen for CPU operations and NVIDIA's cuDNN for optimized GPU performance.



Introduction to Keras

1. Keras, Tensorflow, Theano, CNTK:

➤ Links:

✓ Theano

<http://deeplearning.net/software/theano>

✓ TensorFlow

www.tensorflow.org

✓ CNTK

<https://github.com/Microsoft/CNTK>

✓ Eigen

<http://eigen.tuxfamily.org>



Introduction to Keras

2. Quick overview of Keras:

- ❑ The typical Keras workflow looks just like that example:
 1. **Define** your training data: **input** tensors and target tensors.
 2. **Define** a network of **layers** (or model) that maps your inputs to your targets.
 3. **Configure** the learning process by choosing a **loss** function, an **optimizer**, and some **metrics** to monitor.
 4. **Iterate** on your training data by calling the **fit()** method of your model.



Introduction to Keras

2. Quick overview of Keras:

- ❖ Keras supports two implementation models:

1. Sequential

- ✓ A set of linear layers in the form of a stack.
- ✓ It requires a clear definition of the input data type for it.
- ✓ Easy to learn

2. Functional

- ✓ A model that uses API functions
- ✓ Similar to TensorFlow 2.0

- Reference: <https://keras.io>.



Sequential() model

- ❖ Sequential model and deep learning: same as TensorFlow
 - ✓ Sequential models are suitable for a stack of simple layers where each layer has exactly one input tensor and one output tensor.
 - ✓ Steps to follow
 - Create model
 - Compile model (`model.compile()`)
 - Fit model (`model.fit()`)
 - Model-based prediction (`model.predict()`)
 - Model Evaluation (`model.evaluate()`)
 - Save model (`model.save()`)
 - Load model (`model.load()`)



Sequential() model

- ❖ **Sequential model** and deep learning:
 - ✓ Declare the library packages to use

```
import tensorflow as tf
import keras
from keras import layers
```

Sequential() model

- ❖ Sequential model and deep learning:
 - Create model
 - ✓ Define sequential model have 3 layer

```
model = keras.Sequential(  
    [  
        layers.Dense(2, activation="relu", name="layer1"),  
        layers.Dense(3, activation="relu", name="layer2"),  
        layers.Dense(4, name="layer3"),  
    ]  
)  
  
# Gọi model với một dữ liệu đầu vào  
x = tf.ones((3, 3))  
y = model(x)
```



Sequential() model

- ❖ Sequential model and deep learning:
 - Create model
 - ✓ Another way to declare

```
layer1 = layers.Dense(2, activation="relu", name="layer1")
layer2 = layers.Dense(3, activation="relu", name="layer2")
layer3 = layers.Dense(4, name="layer3")

x = tf.ones((3, 3))
y = layer3(layer2(layer1(x)))
```

Sequential() model

- ❖ Sequential model and deep learning:
 - Create model
 - ✓ Result of executing model creation using method `summary()`

```
Model: "sequential"
-----
Layer (type)                 Output Shape              Param #
-----
layer1 (Dense)               (3, 2)                    8
layer2 (Dense)               (3, 3)                    9
layer3 (Dense)               (3, 4)                    16
-----
Total params: 33
Trainable params: 33
Non-trainable params: 0
-----
```



Sequential() model

- ❖ Sequential model and deep learning:
 - ✓ Compile model (`model.compile()`)

```
model.compile(optimizer='rmsprop',  
              loss='binary_crossentropy',  
              metrics=['accuracy'])
```



Sequential() model

- ❖ Sequential model and deep learning:
 - ✓ Fit model (`model.fit()`)

```
model.fit(X_train, y_train, batch_size=batch_size,  
          epochs=epoch, verbose=1,  
          validation_data=(X_test, y_test),  
          callbacks=[model_checkpoint_callback])
```



Sequential() model

- ❖ Sequential model and deep learning:
 - ✓ Model-based prediction (`model.predict()`)

```
pickle.dump(model, open(file_name, 'wb'))  
  
# Dự đoán lớp trên dữ liệu test  
predicted = model.predict(X_test)
```




Sequential() model

- ❖ Sequential model and deep learning:
 - ✓ Model Evaluation (`model.evaluate()`)

```
scores = model.evaluate(X_test, y_test, verbose=1)
print("Loss:", (scores[0]))
print("Accuracy:", (scores[1]*100))
```



Sequential() model

❖ Sequential model and deep learning

➤ Save model (`model.save()`)

- ✓ Save the model to use in the next call without retraining

```
model.save_weights(model_h5_file)
```

➤ Load model (`model.load()`)

- ✓ Read pre-trained model to use

```
model_temp = model.load_weights(model_h5_file)
```



Keras functional APIs

- ❖ Keras APIs are used to create more flexible models than `keras.Sequential` API.
- ❖ APIs can handle models with non-linear topologies, shared layers, and even multiple inputs or outputs.
- ❖ The main idea behind a deep learning model is a directed acyclic graph (DAG) of layers. So APIs are a way to build a graph of layers.



Keras functional APIs

❖ Example:

- a. Building the model: defining **inputs** and **outputs**

```
inputs = keras.Input(shape=(784,))
img_inputs = keras.Input(shape=(32, 32, 3))
dense = layers.Dense(64, activation="relu")
x = dense(inputs)
x = layers.Dense(64, activation="relu")(x)
outputs = layers.Dense(10)(x)
model = keras.Model(inputs=inputs, outputs=outputs, name="mnist_model")
model.summary()
```

Keras functional APIs

- ❖ Example:
- Results of modeling

Model: "mnist_model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 784)]	0
dense (Dense)	(None, 64)	50240
dense_1 (Dense)	(None, 64)	4160
dense_2 (Dense)	(None, 10)	650

Total params: 55,050

Trainable params: 55,050

Non-trainable params: 0



Keras functional APIs

❖ Example:

b. Training, Evaluation, and Inference

- ✓ The training, evaluation, and inference processes work similarly to **Sequential**.
- ✓ The **Model** class provides a training loop using the **fit()** method.
- ✓ During training, the **Model** performs validation.
- ✓ The process of loading **MNIST** image data involves reshaping the data into vectors, fitting the model to the data.
- ✓ Evaluate the performance of the **Model** on test data using the **evaluate()** method.



Keras functional APIs

- ❖ Example:
 - b. Training, Evaluation, and Inference

```
(x_train, y_train), (x_test, y_test) = tf.keras.datasets.mnist.load_data()
x_train = x_train.reshape(60000, 784).astype("float32") / 255
x_test = x_test.reshape(10000, 784).astype("float32") / 255

model.compile(
    loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
    optimizer=tf.keras.optimizers.RMSprop(),
    metrics=[tf.keras.metrics.SparseCategoricalAccuracy()],
)

history = model.fit(x_train, y_train, batch_size=64, epochs=2, validation_split=0.2)
test_scores = model.evaluate(x_test, y_test, verbose=2)
print("Test loss:", test_scores[0])
print("Test accuracy:", test_scores[1])
```

Keras functional APIs

- ❖ Example:
 - b. **Training**, **evaluation**, and **inference**

- The corpus used in the example is **MINST** (Modified National Institute of Standards and Technology).
- **MINST** contains handwritten digits that are commonly used in training image processing systems.





Keras functional APIs

- ❖ Example:
 - b. Training, Evaluation and Inference
 - Implementation results:

```
Epoch 1/2
750/750 [=====] - 3s 3ms/step - loss: 0.3431 - sparse_ca
Epoch 2/2
750/750 [=====] - 2s 3ms/step - loss: 0.1582 - sparse_ca
313/313 - 0s - loss: 0.1514 - sparse_categorical_accuracy: 0.9565 - 298ms/epoch -
Test loss: 0.15140439569950104
Test accuracy: 0.9564999938011169
```



Keras functional APIs

❖ Example:

c. Save model

- Saving the model is similar to [Sequential](#).
- Call the **save()** method to save the entire model as a file.
- Saving the model as a file allows the same model to be recreated from this file, even if the model generation code is no longer available.
- This saved file contains:
 - ✓ The model architecture
 - ✓ The model weights (obtained from training)
 - ✓ The model training configuration, if any (passed to compilation)
 - ✓ The optimizer and its state (to restart training where it left off)



Keras functional APIs

❖ Example:

c. Save model

```
model.save("path_to_my_model.keras")  
del model  
  
model = keras.models.load_model("path_to_my_model.keras")
```